

## Final Project Submission

Please fill out:

- Student name: Dhruv Ragunathan
- Student pace: part time
- Scheduled project review date/time:
- Instructor name: Mark Barbour
- Blog post URL:

## Introduction and Business Case

As the world struggles to vaccinate the global population against COVID-19, an understanding of how people's backgrounds, opinions, and health behaviors are related to their personal vaccination patterns can provide guidance for future public health efforts. The stakeholders for this study are public health officials responsible for determining vaccination strategy.

The goal of this project is too predict whether people got H1N1 and seasonal flu vaccines using data collected in the National 2009 H1N1 Flu Survey. This is a binary classification problem where we will be investigating if a respondent received the Seasonal flu vaccine.



In this study, we will predict whether a participant will get the seasonal flu vaccine. We will optimize on identifying the most amount of people who need a vaccine, at the expense of erroneously identifying people who have already got the vaccine. In data science terms, we will optimize on reducing the amount of false negatives.

The rationale here is that most people who die from the flu are unvaccinated. Ensuring that people are vaccinated will save lives.

## Data Understanding

- There are three sets of CSVs provided as part of this study.
- Two of the CSVs are for modeling training, and one is for model testing.
- A train - test split has already been done for us, but we will perform a train - test split on the training data anyway.
- The CSV 'training set features' provides columns we can use to target whether a participant got the Flu Vaccine or not
- The CSV training set labels contains the target column.

## Steps to Get the Dataset

Here are the steps to replicate the data

- Go to <https://www.kaggle.com/datasets/imdevskp/h1n1-swine-flu-2009-pandemic-dataset> (<https://www.kaggle.com/datasets/imdevskp/h1n1-swine-flu-2009-pandemic-dataset>)
- Download data.csv
- The data is already separated into a training and testing set
- Features of the data described in training\_set\_labels.csv

## Training Data Set

\* The training data set with the features contain 26,707 entries with 35 different columns. The columns are listed out below.

## Target Training Data Set

- The target feature set contains 26707 entries with 3 columns:
  - Respondent ID
  - H1N1 Vaccine Status
  - Seasonal Flu Vaccine Status

For this study, we are predicting whether the H1N1 Vaccine was used. The 0s in the H1N1 column indicate that the vaccine was not taken. 1 indicates that it was.

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 # Surpress warning errors
        2
        3 import warnings
        4 warnings.filterwarnings('ignore')
```

```
In [3]: 1 X = pd.read_csv('training_set_features.csv', index_col='respondent_id')
        2 y = pd.read_csv('training_set_labels.csv', index_col='respondent_id') #1
```

In [4]: 1 X.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 26707 entries, 0 to 26706
Data columns (total 35 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   h1n1_concern                        26615 non-null  float64
 1   h1n1_knowledge                      26591 non-null  float64
 2   behavioral_antiviral_meds           26636 non-null  float64
 3   behavioral_avoidance                26499 non-null  float64
 4   behavioral_face_mask                26688 non-null  float64
 5   behavioral_wash_hands               26665 non-null  float64
 6   behavioral_large_gatherings         26620 non-null  float64
 7   behavioral_outside_home             26625 non-null  float64
 8   behavioral_touch_face               26579 non-null  float64
 9   doctor_recc_h1n1                   24547 non-null  float64
10   doctor_recc_seasonal                24547 non-null  float64
11   chronic_med_condition               25736 non-null  float64
12   child_under_6_months                25887 non-null  float64
13   health_worker                       25903 non-null  float64
14   health_insurance                    14433 non-null  float64
15   opinion_h1n1_vacc_effective          26316 non-null  float64
16   opinion_h1n1_risk                    26319 non-null  float64
17   opinion_h1n1_sick_from_vacc          26312 non-null  float64
18   opinion_seas_vacc_effective          26245 non-null  float64
19   opinion_seas_risk                    26193 non-null  float64
20   opinion_seas_sick_from_vacc          26170 non-null  float64
21   age_group                           26707 non-null  object
22   education                           25300 non-null  object
23   race                                26707 non-null  object
24   sex                                  26707 non-null  object
25   income_poverty                      22284 non-null  object
26   marital_status                      25299 non-null  object
27   rent_or_own                         24665 non-null  object
28   employment_status                  25244 non-null  object
29   hhs_geo_region                      26707 non-null  object
30   census_msa                          26707 non-null  object
31   household_adults                   26458 non-null  float64
32   household_children                 26458 non-null  float64
33   employment_industry                13377 non-null  object
34   employment_occupation              13237 non-null  object
dtypes: float64(23), object(12)
memory usage: 7.3+ MB

```

In [5]: 1 y.head(5)

Out[5]:

	h1n1_vaccine	seasonal_vaccine
respondent_id		
0	0	0
1	0	1
2	0	0
3	0	1
4	0	0

In [6]: 1 y.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 26707 entries, 0 to 26706
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   h1n1_vaccine     26707 non-null  int64
1   seasonal_vaccine 26707 non-null  int64
dtypes: int64(2)
memory usage: 625.9 KB
```

In [7]: 1 y["h1n1\_vaccine"].value\_counts()

Out[7]: 0 21033  
1 5674  
Name: h1n1\_vaccine, dtype: int64

In [8]: 1 y["seasonal\_vaccine"].value\_counts()

Out[8]: 0 14272  
1 12435  
Name: seasonal\_vaccine, dtype: int64

**Since there is a better mix of data for the seasonal vaccine, we are dropping the h1n1 vaccine column**

In [9]: 1 y.drop(columns = ["h1n1\_vaccine"],axis=0,inplace=True)

```
In [10]: 1 y
```

```
Out[10]:
```

seasonal_vaccine	
respondent_id	
0	0
1	1
2	0
3	1
4	0
...	...
26702	0
26703	0
26704	1
26705	0
26706	0

26707 rows × 1 columns

## Exploratory Data Analysis

Created several slices of the data to explore multiple relationships


- First created a graph comparing the vaccination status of patients based on age
- Second compared education level against vaccination status
- Compared race to vaccination status
- Compared gender to vaccination status

```
In [11]: 1 # Merging data columns to perform exploratory data analysis
          2
          3 seas_df = pd.merge(X,y,on = y.index)
```

```
In [12]: 1 # Compare Age against Vaccination Status
```

```
In [13]: 1 # Creating a histogram of the proportion of participants who were vacci
2
3 sns.countplot(y=seas_df['age_group'],hue=seas_df['seasonal_vaccine'],da
4 sns.set(rc={'figure.figsize':(15,15)})
5 plt.xticks()
6 plt.title('Seasonal Flu Vaccination Status by Age')
7 plt.legend(loc=1)
```


Out[13]: <matplotlib.legend.Legend at 0x7fd0f7f4d400>



```
In [14]: 1 # Compare Education Level Against Vaccination Status
```

```
In [15]: 1 sns.countplot(y=seas_df['education'],hue=seas_df['seasonal_vaccine'],da
2         sns.set(rc={'figure.figsize':(15,15)})
3         plt.xticks()
4         plt.title('Seasonal Flu Vaccination Status by Education Level')
5         plt.legend(loc=1)
```

Out[15]: <matplotlib.legend.Legend at 0x7fd0f6fe1130>




```
In [16]: 1 # Compare Race to Seasonal Vaccination Status
```



```
In [17]: 1 sns.countplot(y=seas_df['race'],hue=seas_df['seasonal_vaccine'],data=seas_df)
2         sns.set(rc={'figure.figsize':(15,15)})
3         plt.xticks()
4         plt.title('Seasonal Flu Vaccination Status by Race')
5         plt.legend(loc=1)
```


Out[17]: <matplotlib.legend.Legend at 0x7fd0f8c31220>



```
In [18]: 1 # Compare Gender to Vaccination Status
```


```
In [19]: 1 sns.countplot(y=seas_df['sex'],hue=seas_df['seasonal_vaccine'],data=seas_df)
2         sns.set(rc={'figure.figsize':(15,15)})
3         plt.xticks()
4         plt.title('Seasonal Flu Vaccination Status by Gender')
5         plt.legend(loc=1)
```

Out[19]: <matplotlib.legend.Legend at 0x7fd0f903b430>



```
In [20]: 1 sns.barplot(x = seas_df['seasonal_vaccine'], y = seas_df['doctor_recc_s
2 plt.xticks()
3 plt.title("Seasonal Flu Vaccination Status based on Doctor's Reccomenda
4 plt.legend(loc=1)
5 plt.show()
```

No handles with labels found to put in legend.



## Data Preparation

Here is the process we followed for data preparation

- Dropped columns not related to Seasonal: 'h1n1\_concern', 'h1n1\_knowledge', 'doctor\_recc\_h1n1', 'opinion\_h1n1\_vacc\_effective', 'opinion\_h1n1\_risk', 'opinion\_h1n1\_sick\_from\_vacc',
- Drop Columns: health insurance, employment\_industry, employment occupation, income\_poverty

- drop rows: h1n1\_, *behavioral*\_, opinion\_h1n1\_vacc\_effective, opinion\_h1n1\_risk, opinion\_h1n1\_sick\_from\_vacc, household\_adults, household\_children
- fill with median: doctor\_recc\_\*, health\_worker
- fill with mode: chronic\_med\_cond, child\_under\_6\_months, health\_worker, education, marital status, rent or own, employment status

```
In [21]: 1 X.isna().sum()
```

```
Out[21]: h1n1_concern          92
h1n1_knowledge          116
behavioral_antiviral_meds    71
behavioral_avoidance        208
behavioral_face_mask         19
behavioral_wash_hands        42
behavioral_large_gatherings   87
behavioral_outside_home      82
behavioral_touch_face        128
doctor_recc_h1n1           2160
doctor_recc_seasonal        2160
chronic_med_condition        971
child_under_6_months         820
health_worker                804
health_insurance            12274
opinion_h1n1_vacc_effective   391
opinion_h1n1_risk            388
opinion_h1n1_sick_from_vacc   395
opinion_seas_vacc_effective   462
opinion_seas_risk            514
opinion_seas_sick_from_vacc   537
age_group                    0
education                    1407
race                          0
sex                           0
income_poverty               4423
marital_status               1408
rent_or_own                  2042
employment_status            1463
hhs_geo_region                0
census_msa                    0
household_adults              249
household_children            249
employment_industry           13330
employment_occupation         13470
dtype: int64
```

In [22]: 1 X.columns

Out[22]: Index(['h1n1\_concern', 'h1n1\_knowledge', 'behavioral\_antiviral\_meds',  
'behavioral\_avoidance', 'behavioral\_face\_mask', 'behavioral\_wash\_h  
ands',  
'behavioral\_large\_gatherings', 'behavioral\_outside\_home',  
'behavioral\_touch\_face', 'doctor\_recc\_h1n1', 'doctor\_recc\_seasona  
l',  
'chronic\_med\_condition', 'child\_under\_6\_months', 'health\_worker',  
'health\_insurance', 'opinion\_h1n1\_vacc\_effective', 'opinion\_h1n1\_r  
isk',  
'opinion\_h1n1\_sick\_from\_vacc', 'opinion\_seas\_vacc\_effective',  
'opinion\_seas\_risk', 'opinion\_seas\_sick\_from\_vacc', 'age\_group',  
'education', 'race', 'sex', 'income\_poverty', 'marital\_status',  
'rent\_or\_own', 'employment\_status', 'hhs\_geo\_region', 'census\_ms  
a',  
'household\_adults', 'household\_children', 'employment\_industry',  
'employment\_occupation'],  
dtype='object')

In [23]: 1 *#drop columns not related to Seasonal: 'h1n1\_concern', 'h1n1\_knowledge'*  
2 *# Drop Columns: health insurance, employment\_industry, employment occup*  
3 *# drop rows: h1n1\_\*, behavioral\_\*, opinion\_h1n1\_vacc\_effective, opinion*  
4 *#household\_adults, household\_children*  
5  
6 *# fill with median: doctor\_recc\_\*, health\_worker*  
7 *# fill with mode: chronic\_med\_cond, child\_under\_6\_months, health\_worker*  
8  
9  
10

In [24]: 1 *# List of columns dropped related to seasonal vaccine*  
2  
3 H1N1\_Columns\_Dropped = ['h1n1\_concern', 'h1n1\_knowledge', 'doctor\_recc\_h  
4

In [25]: 1 *# List of Columns Dropped because they have large amounts of data missi*  
2  
3 Column\_Dropped\_High\_Data\_Loss = ["health\_insurance", "employment\_indust

In [26]: 1 *st of Columns where we are dropping columns with NA*  
2  
3 Row\_Columns = ['behavioral\_avoidance', 'behavioral\_face\_mask', 'behavioral  
4 'behavioral\_large\_gatherings', 'behavioral\_outside\_home', 'behavioral\_tou  
5 'opinion\_seas\_sick\_from\_vacc', 'household\_children', 'household\_adults']  
6

```
In [27]: 1 # List of columns where we fill missing values with the median and Mode
2
3 Fill_Median = ['health_worker']
4 Fill_Mode = ['chronic_med_condition', 'child_under_6_months', 'health_w
5
```

```
In [28]: 1 # First drop seasonal columns
2
3 X = X.drop(H1N1_Columns_Dropped,axis = 1)
4
5 X.columns
```

```
Out[28]: Index(['behavioral_antiviral_meds', 'behavioral_avoidance',
'behavioral_face_mask', 'behavioral_wash_hands',
'behavioral_large_gatherings', 'behavioral_outside_home',
'behavioral_touch_face', 'doctor_recc_seasonal',
'chronic_med_condition', 'child_under_6_months', 'health_worker',
'health_insurance', 'opinion_seas_vacc_effective', 'opinion_seas_r
isk',
'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race',
'sex',
'income_poverty', 'marital_status', 'rent_or_own', 'employment_sta
tus',
'hhs_geo_region', 'census_msa', 'household_adults',
'household_children', 'employment_industry', 'employment_occupatio
n'],
dtype='object')
```

```
In [29]: 1 # Drop columns that have mostly NA rows
2
3 X = X.drop(Column_Dropped_High_Data_Loss,axis = 1)
4
5 X.columns
```

```
Out[29]: Index(['behavioral_antiviral_meds', 'behavioral_avoidance',
'behavioral_face_mask', 'behavioral_wash_hands',
'behavioral_large_gatherings', 'behavioral_outside_home',
'behavioral_touch_face', 'doctor_recc_seasonal',
'chronic_med_condition', 'child_under_6_months', 'health_worker',
'opinion_seas_vacc_effective', 'opinion_seas_risk',
'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race',
'sex',
'marital_status', 'rent_or_own', 'employment_status', 'hhs_geo_reg
ion',
'census_msa', 'household_adults', 'household_children'],
dtype='object')
```

```
In [30]: 1 # Fill columns with median
2
3 for col in Fill_Median:
4     X[col].fillna(X[col].median(),inplace = True)
5
6 X[Fill_Median].isna().sum()
```

```
Out[30]: health_worker      0
dtype: int64
```

```
In [31]: 1 # Fill columns with Mode
2
3 for col in Fill_Mode:
4     X[col].fillna(X[col].mode()[0],inplace=True)
5
6 X[Fill_Mode].isna().sum()
```

```
Out[31]: chronic_med_condition      0
child_under_6_months      0
health_worker      0
education      0
marital_status      0
rent_or_own      0
employment_status      0
dtype: int64
```

```
In [32]: 1 # Finally drop rows in columns that have minimal data missing
2
3 X.dropna(axis=0,inplace=True)
4
5 X[Drop_Row_Columns].isna().sum()
```

```
Out[32]: behavioral_avoidance      0
behavioral_face_mask      0
behavioral_wash_hands      0
behavioral_large_gatherings      0
behavioral_outside_home      0
behavioral_touch_face      0
opinion_seas_vacc_effective      0
opinion_seas_risk      0
opinion_seas_sick_from_vacc      0
household_children      0
household_adults      0
dtype: int64
```

```
In [33]: 1 X.isna().sum()
```

```
Out[33]: behavioral_antiviral_meds      0
behavioral_avoidance                  0
behavioral_face_mask                  0
behavioral_wash_hands                 0
behavioral_large_gatherings           0
behavioral_outside_home               0
behavioral_touch_face                 0
doctor_recc_seasonal                  0
chronic_med_condition                 0
child_under_6_months                 0
health_worker                         0
opinion_seas_vacc_effective           0
opinion_seas_risk                     0
opinion_seas_sick_from_vacc           0
age_group                             0
education                             0
race                                  0
sex                                    0
marital_status                        0
rent_or_own                           0
employment_status                     0
hhs_geo_region                        0
census_msa                            0
household_adults                      0
household_children                    0
dtype: int64
```



```
In [34]: 1 X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 23574 entries, 0 to 26706
Data columns (total 25 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   behavioral_antiviral_meds             23574 non-null  float64
 1   behavioral_avoidance                  23574 non-null  float64
 2   behavioral_face_mask                  23574 non-null  float64
 3   behavioral_wash_hands                 23574 non-null  float64
 4   behavioral_large_gatherings           23574 non-null  float64
 5   behavioral_outside_home               23574 non-null  float64
 6   behavioral_touch_face                 23574 non-null  float64
 7   doctor_recc_seasonal                  23574 non-null  float64
 8   chronic_med_condition                 23574 non-null  float64
 9   child_under_6_months                 23574 non-null  float64
10   health_worker                        23574 non-null  float64
11   opinion_seas_vacc_effective            23574 non-null  float64
12   opinion_seas_risk                      23574 non-null  float64
13   opinion_seas_sick_from_vacc            23574 non-null  float64
14   age_group                             23574 non-null  object
15   education                             23574 non-null  object
16   race                                  23574 non-null  object
17   sex                                   23574 non-null  object
18   marital_status                       23574 non-null  object
19   rent_or_own                          23574 non-null  object
20   employment_status                    23574 non-null  object
21   hhs_geo_region                       23574 non-null  object
22   census_msa                           23574 non-null  object
23   household_adults                     23574 non-null  float64
24   household_children                   23574 non-null  float64
dtypes: float64(16), object(9)
memory usage: 4.7+ MB
```

## Data Preparation Target Dataframe

- Now let's prepare the data in y. In X we removed a lot of rows with respondent IDs that are still in our target.
- let's ensure X and y have the same respondents in the data set.
- We can do this by inner joining X and y, then only selecting for the "seasonal\_vaccine" column

```
In [35]: 1 # Currently there are around 26K rows
          2
          3 y.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 26707 entries, 0 to 26706
Data columns (total 1 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   seasonal_vaccine      26707 non-null  int64
dtypes: int64(1)
memory usage: 417.3 KB
```

```
In [36]: 1 # Only leave rows that are in X
          2
          3 y = pd.concat(
          4     [y,X],
          5     axis=1,
          6     join="inner"
          7 )["seasonal_vaccine"]
          8
          9 y
```

Out[36]:

seasonal_vaccine	
respondent_id	
0	0
1	1
3	1
4	0
5	0
...	...
26701	0
26702	0
26703	0
26704	1
26706	0

23574 rows × 1 columns

```
In [37]: 1 # y now has the same number of rows as X
          2
          3 assert len(y) == len(X)
```

## Modeling

- First import libraries

- Train - Test Split the data
- Start with the baseline logistical model

```
In [38]: 1 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
2 from sklearn.preprocessing import StandardScaler, OneHotEncoder, FunctionTransformer
3 from sklearn.impute import SimpleImputer
4 from sklearn.compose import ColumnTransformer
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.svm import SVC
7 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
8 from sklearn.svm import LinearSVC
9 from sklearn.tree import DecisionTreeClassifier
10 from sklearn.naive_bayes import GaussianNB
11 from sklearn.neighbors import KNeighborsClassifier
12 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, plot_confusion_matrix
13 from sklearn.multioutput import MultiOutputClassifier
14 from sklearn.compose import ColumnTransformer
15 from sklearn.pipeline import Pipeline
```

```
In [39]: 1 ## Train Test Split -
2
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2,
4
```

## Baseline model

- A logistical Regression model will be used as the baseline model.
- Two forms of the model were evaluated. One that used OHE and one that did not.
- Data scaled using a standard scaler after these methods.
- Train-test split occurs after transformation methods to prevent data lost.

```
In [40]: 1 X_train.columns
```

```
Out[40]: Index(['behavioral_antiviral_meds', 'behavioral_avoidance',
'behavioral_face_mask', 'behavioral_wash_hands',
'behavioral_large_gatherings', 'behavioral_outside_home',
'behavioral_touch_face', 'doctor_recc_seasonal',
'chronic_med_condition', 'child_under_6_months', 'health_worker',
'opinion_seas_vacc_effective', 'opinion_seas_risk',
'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race',
'sex',
'marital_status', 'rent_or_own', 'employment_status', 'hhs_geo_reg
ion',
'census_msa', 'household_adults', 'household_children'],
dtype='object')
```

```
In [41]: 1 numeric_columns = list(X_train.columns[X_train.dtypes == 'float64'].value
2 object_columns = list(X_train.columns[X_train.dtypes == 'object'].value)
```

```
In [42]: 1 X_train_object = X_train[object_columns]
2 X_test_object = X_test[object_columns]
3
4 ohe = OneHotEncoder(categories="auto", handle_unknown="ignore", sparse=True)
5
6 X_train_ohe = pd.DataFrame(ohe.fit_transform(X_train_object), columns=ohe.get_feature_names_out())
7 X_test_ohe = pd.DataFrame(ohe.transform(X_test_object), columns=ohe.get_feature_names_out())
8
9 X_train_ohe = pd.concat([X_train[numeric_columns], X_train_ohe], axis=1)
10 X_test_ohe = pd.concat([X_test[numeric_columns], X_test_ohe], axis=1)
```

```
In [43]: 1 # Instantiate a LogisticRegression with random_state=42
2 log_reg = LogisticRegression(solver="liblinear")
3
4 baseline_model = log_reg.fit(X_train_ohe, y_train)
5
6 # Make predictions on test data
7 y_pred = baseline_model.predict(X_test_ohe)
8
9 # Calculate the ROC-AUC score
10 roc_auc = roc_auc_score(y_test, y_pred)
11 print("ROC-AUC score: ", roc_auc)
```

ROC-AUC score: 0.7789120090838635

```
In [44]: 1 ## Compare against model that does not use the OHE method
```

```
In [45]: 1 log_reg = LogisticRegression(solver="liblinear")
2
3 baseline_model = log_reg.fit(X_train[numeric_columns], y_train)
4
5 # Make predictions on test data
6 y_pred = baseline_model.predict(X_test[numeric_columns])
7
8 # Calculate the ROC-AUC score
9 roc_auc = roc_auc_score(y_test, y_pred)
10 print("ROC-AUC score: ", roc_auc)
```

ROC-AUC score: 0.756862462375863

```
In [46]: 1 ## Our model did slightly better without the categorical columns since
```

```
In [47]: 1 ## Standardization and Scaling
```

```
In [48]: 1 #scaling the Dataset
2
3 ss = StandardScaler()
4
5 X_scaled = ss.fit_transform(X[numeric_columns])
6 X_scaled = pd.DataFrame(X_scaled, columns=numeric_columns)
7 X_scaled
```

Out[48]:

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_was
respondent_id				
0	0.0	0.0	0.0	
1	0.0	1.0	0.0	
3	0.0	1.0	0.0	
4	0.0	1.0	0.0	
5	0.0	1.0	0.0	
...	...	...	...	
26701	0.0	0.0	0.0	
26702	0.0	1.0	0.0	
26703	0.0	1.0	0.0	
26704	0.0	1.0	1.0	
26706	0.0	1.0	0.0	

23574 rows × 16 columns

```
In [49]: 1 # Since we've lost rows from scaling the data, we need to create y_scaled
2
3 y_scaled = pd.concat(
4     [y, X_scaled[numeric_columns]],
5     axis=1,
6     join="inner"
7 )["seasonal_vaccine"]
8
9 y_scaled.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 23574 entries, 0 to 26706
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   seasonal_vaccine  23574 non-null  int64
dtypes: int64(1)
memory usage: 368.3 KB
```

```
In [50]: 1 assert len(X_scaled) == len(y_scaled) #X_scaled and y_scaled are not th
```

```
In [51]: 1  ## Train Test Split for scaling
          2
          3  X_train_s, X_test_s, y_train_s, y_test_s = train_test_split(X_scaled, y
```

```
In [52]: 1  # Create Baseline model for Seasonal Flu Vaccine
          2  baseline_model.fit(X_train_s, y_train_s)
          3
          4  # Make predictions on test data
          5  y_pred2 = baseline_model.predict(X_test_s)
          6
          7  # Calculate the ROC-AUC score
          8  roc_auc = roc_auc_score(y_test_s, y_pred2)
          9  print("ROC-AUC score: ", roc_auc)
```


ROC-AUC score: 0.756862462375863

```
In [53]: 1  # Scaling does not improve the ROC-AUC score.
```

```
In [54]: 1  # Calculate the Confusion Matrix Values
          2
          3  cm = confusion_matrix(y_test, y_pred2)
          4  TN, FP, FN, TP = confusion_matrix(y_test_s, y_pred2).ravel()
          5
          6  print('True Positive(TP) = ', TP)
          7  print('False Positive(FP) = ', FP)
          8  print('True Negative(TN) = ', TN)
          9  print('False Negative(FN) = ', FN)
```

True Positive(TP) = 1648  
False Positive(FP) = 529  
True Negative(TN) = 1926  
False Negative(FN) = 612

```
In [55]: 1 from sklearn import metrics
2
3 confusion_matrix = metrics.confusion_matrix(y_test_s,y_pred2)
4
5 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix)
6
7 cm_display.plot()
8 plt.show()
```



## Evaluation of Baseline Model

Our baseline model has identified 1648 as true positive results. The model identified the result as positive and its true value was positive.

It had 529 False positive results. Where the model identified them as positive, but these values were actually negative.

It had 612 False Negative results. Where the model identified them as negative, but they were actually positive.

Finally, it identified 1926 true negatives. These values were negative and the model identified them

Since we are trying to predict the status of someone needing a vaccine, it's preferable to reduce the number of False negatives. The use of this model would be to determine who needs a vaccination.

If the model flags a false positive, then a person who got vaccinated, would be pinged again for vaccination.

A false negative would go under the radar, potentially not receive a vaccination, which could be detrimental.

## Evaluation of Other Models

- Let's evaluate other models to reduce the false negative count of the previous model
- The following models were evaluated in addition to Logistic Regression: Naive Bayes, Support Vector Machines, Decision Trees, Random Forest, and K-Nearest Neighbor.
- We calculated the following metrics to evaluate these models against one-another: accuracy, precision, recall, and Roc-AUC score.

```
In [56]: 1  # Initialize the other model classifiers such as logistic regression, r  
2  # random forest, and K-nearest neighbors  
3  
4  models = {}  
5  
6  models['Logistic Regression'] = LogisticRegression()  
7  
8  models['Naive Bayes'] = GaussianNB()  
9  
10 models['Support Vector Machines'] = LinearSVC()  
11  
12 models['Decision Trees'] = DecisionTreeClassifier()  
13  
14 models['Random Forest'] = RandomForestClassifier()  
15  
16 models['K-Nearest Neighbor'] = KNeighborsClassifier()  
17
```



```

In [57]: 1 # loop over each classifier to evaluate poerformance
2 acc, rec, prec, F1, Roc_Auc = {}, {}, {}, {}, {}
3
4 for model in models.keys():
5
6     # Fit the classifier
7     models[model].fit(X_train_s, y_train_s)
8
9     # Make predictions
10    y_pred3 = models[model].predict(X_test_s)
11
12    # Calculate metrics
13    acc[model] = accuracy_score(y_pred3, y_test_s).ravel()
14    rec[model] = recall_score(y_pred3, y_test_s).ravel()
15    prec[model] = precision_score(y_pred3, y_test_s).ravel()
16    F1[model] = f1_score(y_pred3, y_test_s).ravel()
17    Roc_Auc[model] = roc_auc_score(y_test_s, y_pred3)
18
19
20

```

```

In [58]: 1 metrics = pd.DataFrame(index=models.keys(), columns=['Accuracy', 'Recall', 'Precision', 'F1 Score'])
2 metrics['Accuracy'] = acc.values()
3 metrics['Recall'] = rec.values()
4 metrics['Precision'] = prec.values()
5 metrics['F1 Score'] = F1.values()
6 metrics['Roc-AUC Score'] = Roc_Auc.values()
7 metrics

```

Out[58]:

	Accuracy	Recall	Precision	F1 Score
<b>Logistic Regression</b>	(0.7580063626723224,)	(0.7567691601652135,)	(0.7296460176991151,)	(0.7429601261545392,)
<b>Naive Bayes</b>	(0.7257688229056204,)	(0.7117827420061322,)	(0.7190265486725663,)	(0.7153863086066475,)
<b>Support Vector Machines</b>	(0.7571580063626723,)	(0.7584608252202133,)	(0.7238938053097345,)	(0.7407742811863256,)
<b>Decision Trees</b>	(0.6981972428419937,)	(0.6999522216913521,)	(0.6482300884955752,)	(0.6730990121755112,)
<b>Random Forest</b>	(0.7266171792152704,)	(0.7230133210840606,)	(0.6964601769911505,)	(0.7094883930583727,)
<b>K-Nearest Neighbor</b>	(0.7223753976670202,)	(0.7200370198981952,)	(0.6884955752212389,)	(0.7039131418231169,)


```
In [59]: 1 # Plot features in bar graph
2 # Convert objects to strings
3
4 metrics_num = metrics
5
6 for col in metrics:
7     if col != 'Roc-AUC Score':
8         for val in metrics[col]:
9             value = val[0] #Extract the numeric value from the tuple
10             metrics_num[col] = value
11
12 metrics_num
```

Out[59]:

	Accuracy	Recall	Precision	F1 Score	Roc-AUC Score
<b>Logistic Regression</b>	0.722375	0.720037	0.688496	0.703913	0.756880
<b>Naive Bayes</b>	0.722375	0.720037	0.688496	0.703913	0.725501
<b>Support Vector Machines</b>	0.722375	0.720037	0.688496	0.703913	0.755837
<b>Decision Trees</b>	0.722375	0.720037	0.688496	0.703913	0.696213
<b>Random Forest</b>	0.722375	0.720037	0.688496	0.703913	0.725419
<b>K-Nearest Neighbor</b>	0.722375	0.720037	0.688496	0.703913	0.721030

```
In [60]: 1 metrics_num.plot.bar()
```

```
Out[60]: <AxesSubplot:>
```



## Final Model Evaluation

All models have fairly close accuracy, recall, precision, and roc-auc scores. Based on the business problem, we would want to reduce the number of false negatives. A false negative implies that someone was not vaccinated, was predicted to not need a vaccination. That means that this person would be missed if this model would target whom to outreach.

As a result, the logistic regression model and the Support Vector machine model are the top models. Both have accuracy scores of around 76% and recall of 76%.

It seems like the baseline model using logistical regression performed better on the test data with a higher recall and accuracy score. As a result, the final model for this project will be the Logistical regression model.